

Prediction settling across loops (exploratory; frozen read-out after each loop, 2M validation tokens; difficulty = decile of the token's loss at the training r)

