

20M middle block, 32N-token trunk checkpoints: random loop-count training converges to a fixed point
(robust to the loop count, no gain beyond ~ 4 loops); star = nominal $r=8$

