

20M middle block, 32N-token trunk checkpoints; star = nominal $r=8$
random loop-count training converges to a fixed point: robust, but no gain beyond ~ 4 loops

