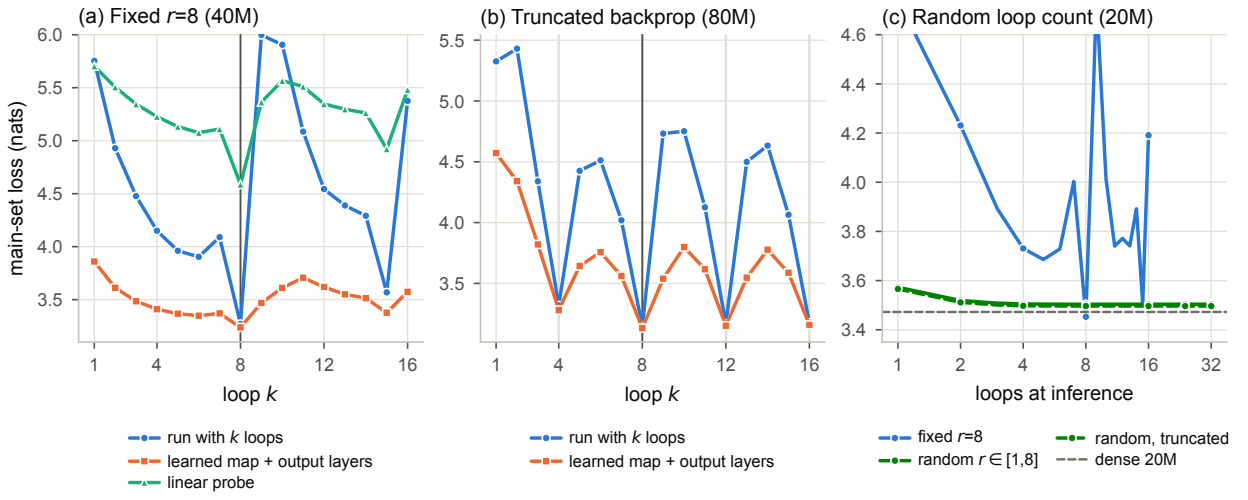
(a) Fixed r=8 (40M)
(b) Truncated backprop (80M)
(c) Random loop count (20M)
main-set loss (nats)
loop k
loop k
loops at inference
run with k loops
learned map + output layers
linear probe
run with k loops
learned map + output layers
fixed r=8
random r ∈ [1,8]
random, truncated
dense 20M